

# Neural Network-Based Character-to-Symbol Sequence Translation for Text Compression

Giuseppe Prisco - 1895709

**“Engineering in Computer Science”** Master’s Degree  
**Deep Learning** (2023/2024)

# Task Description (1)

The goal of the project was to develop a neural network model that could convert **human-readable text** into a **machine-readable symbol sequence** that is compatible with unzip softwares.

- “The dogs are cute” ➡ x\x9c\x0b\xc9HUH\xc9O/VH,JUH.-l\x05\x005\x7f\x06\x18
- “The dogs are cute” ➡ eJwLyUhVSMIPL1ZILEpVSC4tSQUANX8GGA==

# Task Description (2)

Given the **innovative** nature of this project, I explored various approaches related to:

- **pre-processing** and **post-processing** of input and output data formats
- alternative **architectures** to accomplish the task

# Dataset Creation and Pre-processing

## Dataset Creation

Download the “**GLUE**” dataset from the Hugging Face Hub and divide it into:

- **training** set
- **validation** set
- **test** set

## Pre-process Data

- build pairs of **original texts** and their **compressed version**:
  - just compress with zlib
  - compress with zlib and then encode in base64
- **encode** each string of the pair using T5 or ByT5 tokenizer

# Implementation of the Baseline

The tested baseline consist of a simple **Recurrent Neural Network** (RNN) with the following structure:

- **Embedding Layer** ➡ the integers representing tokens pass through this layer, which converts each integer into a dense vector (embeddings)
- **Gated Recurrent Unit (GRU) Layer** ➡ it process the embeddings, maintaining hidden states, and returns the sequence of hidden states for each step
- **Fully Connected Linear Layer** ➡ it projects each hidden state back to the size of the vocabulary, returning the raw predictions of the model

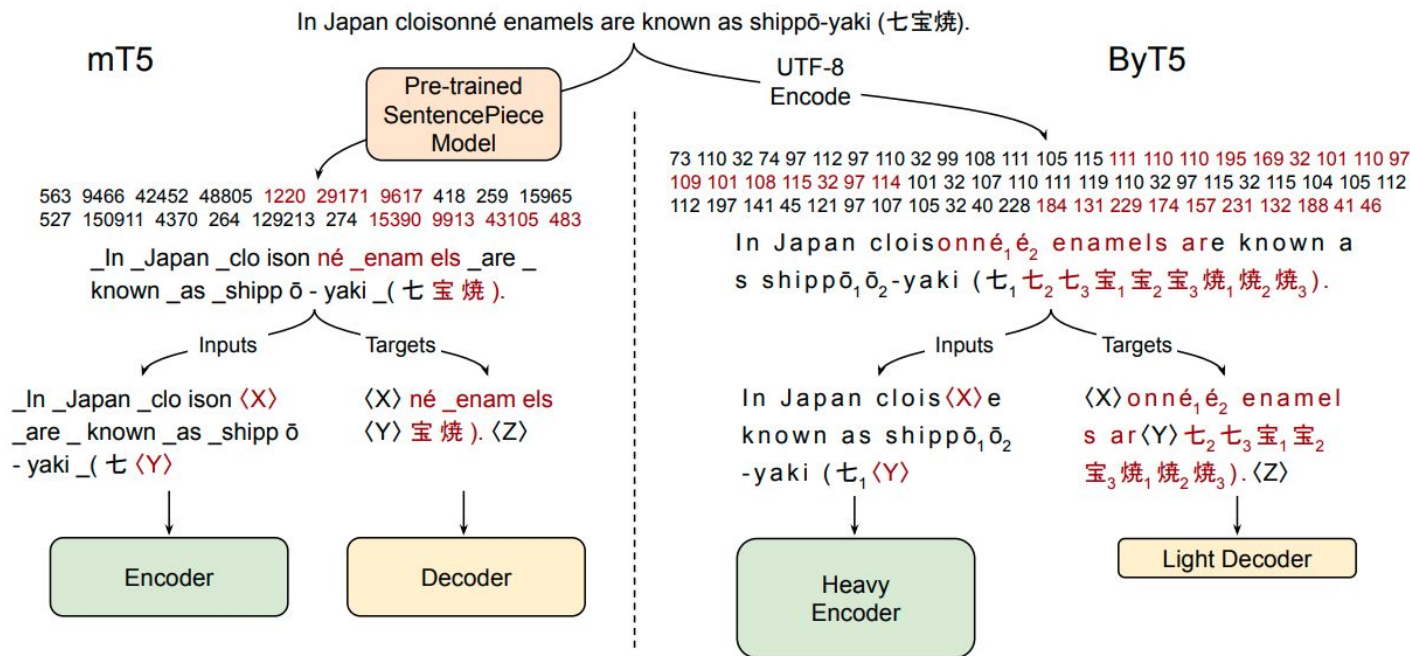
# Final Model Architecture (1)

The final system is based on the T5 model from the Hugging Face Transformers library, specifically **T5ForConditionalGeneration** model.

The two tested models differ on the architecture employed ➔ **t5\_base** or **byt5\_base**:

- **Embedding** layer
- stack of **Encoder** layers
- stack of **Decoder** layers
- **Fully Connected** linear layer

# Final Model Architecture (2)



# Final Model Architecture (3)

The model is trained with the following:

- **"input\_ids"** ➡ they correspond to the tokenized original texts
- **"attention\_mask"** ➡ it differentiates between tokens that are actual words versus others that just represent the padding
- **"labels"** ➡ they correspond to the tokenized compressed texts

The model outputs a **probability vector**, of size equal to the **"vocab\_size"**, for each element of the input sequence and for each sequence in the batch.



# Final Model Architecture (4)

## T5

- max sequence length → 256
- encoder layers → 12
- dropout → 0.1
- learning rate → 0.00001
- batch size → 8

## ByT5

- max sequence length → 64
- encoder layers → 14
- dropout → 0.1
- learning rate → 0.00001
- batch size → 8

# Implemented Metrics

## Levenshtein distance

- It is a **string metric** used to measure the **difference** between two sequences
- It represents the minimum number of **single-character edits** required to change one sequence into another

## Unzip metric

- It represents how many predicted sequences can successfully be **decompressed**
- This metric assigns a **score of 1** if the zlib library can **successfully decompress** the model's output sequence, and **0 if it cannot**

# Auxiliary Metrics

## Predicted Sequence Length

- It measures the length of the sequence **predicted** by the model

## True Sequence Length

- It represents the length of the **ground truth** sequence

## Min Levenshtein distance

- $\text{abs}(\text{predicted length} - \text{true length})$

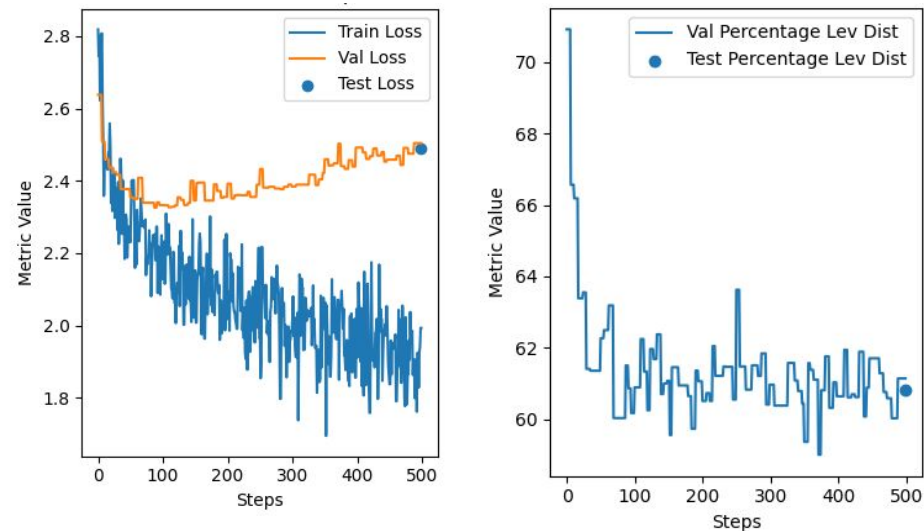
## Max Levenshtein distance

- $\text{max}(\text{predicted length}, \text{true length})$


## % Levenshtein distance

- $((\text{lev} - \text{min}) / (\text{max} - \text{min})) * 100$

# MyBaseline - Results




Validation/Test Losses and Percentage Levenshtein Distance, using byt5, no base64 encoding used

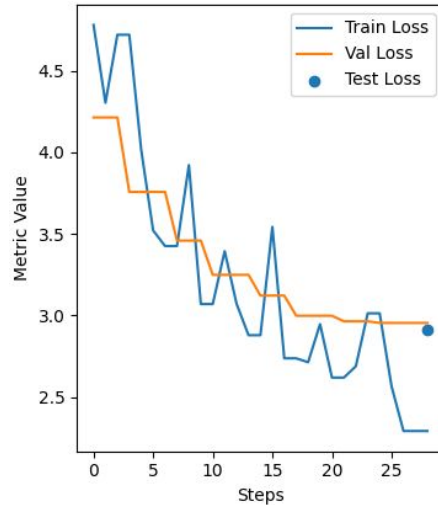



Validation/Test Losses and Percentage Levenshtein Distance, using byt5, base64 encoded

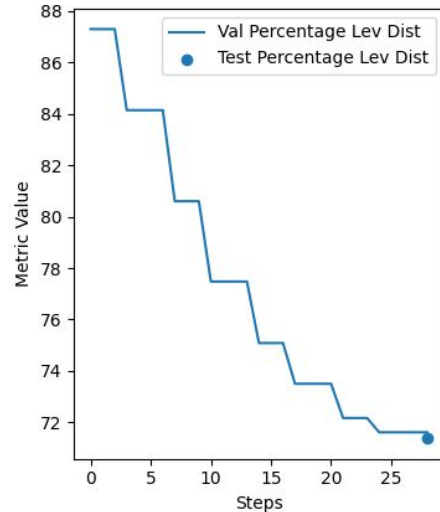
# TextCompressionModel - Results



Validation/Test Losses and Percentage Levenshtein Distance, using t5, no base64 encoding used



Validation/Test Losses and Percentage Levenshtein Distance, using byt5, base64 encoded



# Summary of Results

In this final table are summarized the **best results** achieved under the runtime constraints imposed by Colab:

Table 1: Models Validation/Test scores for the adopted Metrics

Model Names	Levenshtein Distance (in %)		Unzip Metric	
	No Encoding	Base 64	No Encoding	Base 64
MyBaseline (t5)	68.023	87.781	0	0
MyBaseline (byt5)	60.818	74.563	0	0
TextCompressionModel (t5)	37.347	71.688	0	0
TextCompressionModel (byt5)	37.154	71.382	0	0

As we can see, the best results are obtained when **no encoding** is used, with the “TextCompressionModel” able to achieve better results in comparison to “MyBaseline”

# Conclusions

- comparison of “**MyBaseline**” based on GRUs with “**TextCompressionModel**” based on T5
- **great results** achieved for the Percentage Levenshtein Distance metric
- **challenges**: the model would need to achieve a distance of exactly 0% in order to correctly mimic the behaviour of the unzipping tools ➡ unzip metric is always 0
- the model demonstrated its ability to **recognize** the patterns of any compressed string, hinting at future possibilities to successfully perform the given task

In conclusion, the possibility of having a model that could successfully compress and decompress any given text would provide significant efficiency in text compression.

Thanks for the  
Attention